

HIGH PERFORMANCE COMPUTING for SCIENCE & ENGINEERING (HPCSE) I

HS 2021

EXERCISE 05: MONTE CARLO INTEGRATION & OPENMP

Pascal Weber (webepasc@ethz.ch)

Computational Science and Engineering Lab
ETH Zürich

25.11.2021

Outline

- I. Exercise 1 (Monte-Carlo Integration)
- II. Exercise 2/3 (OpenMP Bughunt)

Monte-Carlo Integration

We want to perform integration $I = \int_{\Omega} f(x) \, dx$ over the domain Ω

From probability theory we know that the expectation value of $f(x)$ for a uniform distribution $\mathcal{U}_{\Omega}(x)$ over the domain Ω reads

$$\mathbb{E}_{x \sim \mathcal{U}_{\Omega}(\cdot)}[f(x)] = \frac{1}{|\Omega|} \int_{\Omega} f(x) \, dx$$

From the central-limit theorem we know we can approximate an expectation value by using samples $x_i \sim \mathcal{U}_{\Omega}(\cdot)$ with $i = 1, \dots, N$ from the distribution

$$\mathbb{E}_{x \sim \mathcal{U}_{\Omega}(\cdot)}[f(x)] \approx \frac{1}{N} \sum_{i=1}^N f(x_i)$$

Combining this knowledge is the foundation for **Monte-Carlo Integration**, where we approximate the integral by evaluating the function at randomly sampled locations

$$I \approx \frac{|\Omega|}{N} \sum_{i=1}^N f(x_i)$$

Estimating π using Monte-Carlo Integration

We know that we can write the volume of an object as an integral over the characteristic function χ in an enclosing domain Ω

$$V = \int_{\Omega} \chi(x) \, dx$$

For the unit circle, the characteristic function is

$$\chi(x, y) = \begin{cases} 1, & x^2 + y^2 \leq 1, \\ 0, & \text{otherwise.} \end{cases}$$

We know that the area of the circle is π and thus we know that

$$\pi = 4 \int_{[0,1]^2} \chi(x, y) \, dx dy$$

Using Monte-Carlo integration we can perform the integral by taking samples $x_i \sim \mathcal{U}_{[0,1]}(\cdot)$ and $y_i \sim \mathcal{U}_{[0,1]}(\cdot)$ with $i = 1, \dots, N$

$$\pi \approx \frac{4}{N} \sum_{i=1}^N f(x_i, y_i)$$

[Pseudo] Random Number Generators

It is hard to get true random numbers. In practice we therefore use pseudo random numbers, which are **deterministic sequences**!

Examples:

- Congruential Generators: $x_i = (c \cdot x_{i-1}) \bmod p$
→ maximal period for Mersenne prime numbers $p \equiv M_q = 2^q - 1$ and $c^{p-1} \bmod p = 1$
- Lagged-Fibonacci: $x_i = x_{i-a} \oplus x_{i-b} := (x_{i-a} + x_{i-b}) \bmod 2, \quad a < b$

We will use **std::default_random_engine** to generate pseudo random numbers.

«It is the library implementation's selection of a generator that provides at least acceptable engine behavior for relatively casual, inexperienced, and/or lightweight use.»

The one that is 'normally' used is **std::mt19937** (Mersenne-Twister with sequence length $2^{19937} - 1$, uniform distribution and fast)

Exercise 1 - montecarlo.cpp

Implementation of characteristic function

```
7 // Characteristic function for unit circle
8 inline double F(double x, double y)
9 {
10     if (x * x + y * y < 1.) { // inside unit circle
11         return 1.;
12     }
13     return 0.;
14 }
```

Implementation of Monte-Carlo integration

```
16 // Method 0: serial
17 double C0(size_t N)
18 {
19     // random generator with seed 0
20     std::default_random_engine g(0);
21
22     // uniform distribution in [0, 1]
23     std::uniform_real_distribution<double> u;
24
25     // perform Monte-Carlo integration
26     double pi = 0.;
27     for (size_t i = 0; i < N; ++i) {
28         double x = u(g);
29         double y = u(g);
30         pi += F(x, y);
31     }
32     return 4 * pi / N;
33 }
```

you are asked to parallelize
this code in several ways

**be careful with the random
number generators /
distributions**

```
34 // Method 1: parallelize C0 without using arrays
35 double C1(size_t N)
36 {
37     double pi = 0.;
38
39     // TODO: Implement Method 1
40
41     return 4 * pi / N;
42 }
43
44 // Method 2: parallelize C0 only using
45 // `omp parallel for reduction`, use arrays without padding
46 double C2(size_t N)
47 {
48     double pi = 0.;
49
50     // TODO: Implement Method 2
51
52     return 4 * pi / N;
53 }
54
55 // Method 2: parallelize C0 only using
56 // `omp parallel for reduction`, use arrays without padding
57 double C3(size_t N)
58 {
59     double pi = 0.;
60
61     // TODO: Implement Method 3
62
63     return 4 * pi / N;
64 }
```

**padding should be of the size of one
cache line [usually 64 bytes]**

➡ The goal of the exercise is to observe the
problems that can occur from false-sharing

Exercise 1 - Makefile

Compile and Launch Benchmarks

for OpenMP compatibility. Use
env2lmod; module load gcc/8.2.0

```
1 CXX = g++
2 CXXFLAGS = -O3 -Wall -Wextra -pedantic -std=c++14 -fopenmp
3
4 all: montecarlo
5
6 montecarlo: montecarlo.cpp
7     $(CXX) $< $(CXXFLAGS) -o montecarlo
8
9 run: montecarlo
10     ./varym $(N)
11
12 plot: run
13     ./plot
14
15 clean:
16     rm -rf montecarlo out results.png
17
18 .PHONY: all clean run plot runplot
```

compiles the executable

runs benchmark

plot results

```
1 #!/bin/bash
2
3 gnuplot << EOF
4
5 set terminal pngcairo
6
7 set xlabel 'threads'
8 set ylabel 'time [s]'
9 set output 'results.png'
10 set grid
11 set key Left left bottom
12 set logscale x 2
13 set logscale y 2
14 t0 = `sed 's,.*,,;q' out/m0`
15 set style data lp
16 plot \
17     "out/m0" lw 3 pt 7 t 'm=0, serial', \
18     "out/m1" lw 3 pt 7 t 'm=1, no arrays', \
19     "out/m2" lw 3 pt 7 t 'm=2, no padding', \
20     "out/m3" lw 3 pt 7 t 'm=3, padding', \
21     t0/x w l lw 1 lc 'black' t 'ideal'
22
23 EOF
```

some syntax

\$ to access content of variable
@ contains command-line arguments
> forwards output
&2 is stderr
| pipelining (forward output to next)
commands -> check `man xyz`
set setting of bash
echo print variable
eval runs the expression
grep search words

varym

```
1 #!/bin/bash
2
3 set -eu
4
5 o=out
6
7 rm -f $o/m*
8 mkdir -p $o
9
10 for m in 0 1 2 3 ; do
11     c="./varynt $m $@ > $o/m$m"
12     echo "$c" >&2
13     eval "$c"
14 done
```

varynt

```
1 #!/bin/bash
2
3 set -eu
4
5 maxnt=${OMP_NUM_THREADS:-4}
6
7 for nt in `seq 1 $maxnt` ; do
8     c="OMP_NUM_THREADS=$nt ./montecarlo $@"
9     echo "$c" >&2
10     o=`eval "$c"`
11     wt=`echo "$o" | grep time | cut -d' ' -f2`
12     echo "$nt" "$wt"
13 done
```

Exercise 2

Question 2: OpenMP Bug Hunting I (20 points)

Identify and explain any bugs in the following OpenMP code. Propose a solution. Assume all headers are included correctly.

```
1  #define N 1000
2
3  extern struct data member[N]; // array of structures, defined elsewhere
4  extern int is_good(int i); // returns 1 if member[i] is "good", 0 otherwise
5
6  int good_members[N];
7  int pos = 0;
8
9  void find_good_members()
10 {
11     #pragma omp parallel for
12     for (int i=0; i<N; i++) {
13         if (is_good(i)) {
14             good_members[pos] = i;
15
16             #pragma omp atomic
17             pos++;
18         }
19     }
20 }
```

Hints:

- In your solution you can use "omp critical" or "omp atomic capture"¹

Exercise 3

Question 3: OpenMP Bug Hunting II (20 points)

a) Identify and explain any *bugs* in the following OpenMP code. Propose a solution. Assume all headers are included correctly.

```
1 // assume there are no OpenMP directives inside these two functions
2 void do_work(const float a, const float sum);
3 double new_value(int i);
4
5 void time_loop()
6 {
7     float t = 0;
8     float sum = 0;
9
10    #pragma omp parallel
11    {
12        for (int step=0; step<100; step++)
13        {
14            #pragma omp parallel for nowait
15            for (int i=1; i<n; i++)
16            {
17                b[i-1] = (a[i]+a[i-1])/2.;
18                c[i-1] += a[i];
19            }
20
21            #pragma omp for
22            for (int i=0; i<m; i++)
23                z[i] = sqrt(b[i]+c[i]);
24
25            #pragma omp for reduction(+:sum)
26            for (int i=0; i<m; i++)
27                sum = sum + z[i];
28
29            #pragma omp critical
30            {
31                do_work(t, sum);
32            }
33
34            #pragma omp single
35            {
36                t = new_value(step);
37            }
38        }
39    }
40 }
```

b) Identify and explain any *improvements* that can be made in the following OpenMP code. Propose a solution. Assume all headers are included correctly.

```
1 void work(int i, int j);
2
3 void nesting(int n)
4 {
5     int i, j;
6     #pragma omp parallel
7     {
8         #pragma omp for
9         for (i=0; i<n; i++)
10        {
11            #pragma omp parallel
12            {
13                #pragma omp for
14                for (j=0; j<n; j++)
15                {
16                    work(i, j);
17                }
18            }
19        }
20    }
21 }
```